non-cognitivism seems to come up once more. For instance one may suppose that the principles of practical reason must be self-evident truths known by intuition, and then Kant will come out looking like a traditional rationalistic realist. This is not just a fantasy—this is how Kant was actually read by many late nineteenth and twentieth century philosophers, especially in Britain. Sidgwick and Mill, who read Kant this way, will serve as sufficiently distinguished examples.

Alternatively, someone who takes seriously Kant's thesis that moral laws are the laws of autonomy, legislated by the agent's own will, may read him, as Hare sometimes seems to do, as a prescriptivist and so a non-cognitivist. So although practical reason theories might at first seem to fall between the cracks, there are ways of making them fit the mold.

I will come back to the question how practical reason theories, or at least Kant's theory, should be understood. What I want to do now is articulate a contrast between the theory of normative concepts that I believe stands behind the debate between the cognitivists and the non-cognitivists, and another theory of normative concepts which I take to be a genuine alternative. I call this alternative, in deference to prevailing usage, constructivism. In order to articulate the contrast, I am going to compare arguments from two of the giants of twentieth century moral philosophy, Bernard Williams and John Rawls. Both of these arguments concern a favorite twentieth century theme—the implications of the diversity of ethical opinion. It is going to be necessary for me to spell these arguments out in some detail. But what I am asking you to be interested in for the purposes of this essay is not their particular success, but rather the conception of moral concepts and along with it of moral philosophy that lies behind them.

III. CASE STUDY ONE: BERNARD WILLIAMS[26]

Williams is certainly not a moral realist, so his position here as its spokesman may occasion some surprise. I choose him for two reasons: first because I think his attempt to articulate the idea behind realism is unrivalled in its clarity, and second because of a realist assumption that I think in the end still haunts his own account of moral objectivity. In *Ethics and the Limits of Philosophy*, Williams argues that there is a contrast between the kind of objectivity we can hope to find in science and that which we can hope to achieve in ethics. He frames this contrast in terms of convergence, that is, in terms of what might lead us to the best kind of agreement about the judgments in question. In science, the ideal form of convergence would be this: we come to agree with one another in our scientific judgments because we are all converging on a description of the way the world really is. In ethics, Williams thinks, this sort of convergence is unavailable, and so another must be found.[27]

Williams's account of realism emerges when he explains what he means by "the way the world really is." One thing we might mean in talking about

"the way the world really is" is whether we have applied our concepts correctly. If we say the sky is blue on a day when it is blue then we have done that. But we can also query our conceptual scheme itself. We can ask whether it is the correct one or the best one or something along those lines. Since science leads us to modify our conceptual scheme, and we think of these modifications as improvements, it seems as if some such question must be in order. Williams supposes that the improved scheme is improved in the sense that it comes closer to describing the way the world really is.

Williams proposes that we can capture the notion of "the way the world really is" by formulating a kind of limiting conception which he calls "the 'absolute conception' of the world."[28] The idea involves a contrast between concepts that are more and less dependent on the particular perspective from which we view the world. For instance, we use color concepts because we see in color, so color concepts are dependent on our own particular perspective. The concept of a certain wave length of light is supposed to be less dependent on our perspective.

Williams associates two other properties with a concept's greater independence from particular perspectives. First, our use of concepts which are more dependent on our own perspectives will be both explained, and justified, in terms of a theory that employs concepts which are less dependent. So for instance our use of color concepts might be explained by a theory of vision which employs wave length concepts. Relatedly this theory will also *justify* our belief that color vision is a form of *perception*, that is, a way of learning about the world, by the way that it explains it.[29] Color vision is a way of learning about the world because it gives us information about wave lengths, or something yet more ultimate, which we take to be part of reality. Second, the more independent of our own perspective a concept is, the more likely it is that it could be shared by other rational investigators who were unlike us in their ways of learning about the world. Suppose that there are rational creatures on some other planet who cannot see colors but can hear them. They could not use color concepts but they might be able to use wave-length concepts. The more independent concepts are more shareable.

Williams thinks that the nearest thing we have to a conception of "the way the world really is" is the conception of the world that is maximally independent of our own perspective. And he thinks that if we and the alien investigators began to converge on such a conception (and of course to agree on what judgments are correct within it) then we would have reason to believe we were converging on the absolute conception. This would be the best case of convergence for science. Our theories would come to converge with the theories of other investigators because all of us were converging on a description of the way the world really is.

Now consider what the parallel would be in ethics. Here Williams first deals with a possible objection, namely that there is nothing analogous to

perceptual judgments in ethics. Seeing the facts is one thing, and evaluating them in a certain way is another. To counter this objection, Williams appeals to the existence of what he calls "thick" as opposed to "thin" ethical concepts. Thin ethical concepts—like right and good and ought—do not appear to be world-guided, in the sense that their application does not appear to be guided by the facts. (Notice the echo of verificationism here.) But thick ethical concepts—Williams's examples are coward, lie, brutality, and gratitude—are world-guided and action-guiding at the same time.[30] Only an action that is in a particular way motivated by fear can be called cowardly, and yet to call an action cowardly is to suggest that it ought not to be done.[31]

Of course a prescriptivist or emotivist or a latter-day expressivist has his own account of these concepts. He thinks that their world-guidedness is one thing and that their action-guidingness is another. To say that an action is cowardly is simply to denigrate an action motivated by fear. The difficulty with this analysis, according to Williams, is that it suggests that it would be possible to use a thick ethical concept with perfect accuracy even if you were completely incapable of appreciating the value that it embodies. Williams thinks this is implausible. He does not mean that we can only use an evaluative concept when we ourselves actually endorse the value in question. But he thinks we can apply such concepts only by entering imaginatively into the world of those who have the relevant values, not merely by applying a set of factual criteria.[32] We have to see the world through their eyes. This makes it natural to think of judgments employing thick ethical concepts as a kind of perceptual judgments, for they are a sort of lens through which we view the world. And that in turn makes it natural to think that, like other perceptual judgments, they may be a kind of knowledge.

I say that the sky is blue, and my visitor from Mars says that it makes a humming noise. Are we agreeing? Certainly we don't *mean* the same thing, since I am talking about how the sky looks and he is talking about how it sounds. Yet when we reflect on these views we find that the things we both say have implications expressible in terms of a more absolute concept, that of wave-lengths. And when we look at those implications our judgments are found to converge. Here we find grounds for confidence that both of our perceptions are guiding us rightly: they are ways of knowing about the world.

To get the parallel in ethics, we would have to compare judgments made in alternative sets of thick ethical concepts. One person says that lying is sinful and another says it is dishonorable. They do not appear to *mean* exactly the same thing. But we might take both of their remarks to have implications describable in terms of what *we* think is a more absolute concept—that lying is *wrong*—and here we find that they converge. This would be evidence that their moral perceptions are guiding them rightly, and are ways of knowing about the moral part of the world. On the other hand, suppose what is in question is avenging an insult. The first person thinks it is sinful to avenge it,

while the second thinks it is dishonorable not to. Does this imply any disagreement, in particular a disagreement about whether avenging an insult is, in our sense, morally wrong? Williams does not think so, and this is partly because there is a world-guided side to the idea of something's being dishonorable. The second person may indeed be dishonored—he may be personally diminished, in his own eyes and those of his community—if he does not stand up for himself. Williams thinks that facts of this kind should make us doubt whether the two people are using concepts that converge on what is morally wrong, and so should also make us doubt whether "morally wrong" is a more absolute concept after all—that is, whether it describes some reality.

Instead Williams proposes a different way in which we might think about differences in ethical beliefs. He suggests that we should see the values of a culture not as their best approximations of the truth about right and wrong, but rather as a kind of *habitation*—although, as he emphasizes, and this is a point I will come back to—*not one that they have built*. Their values form a part of the structure of the social world in which they live. This does not mean that we cannot make any evaluative judgments about a culture's values. We can ask whether the social world that is made of those values is a good place for human beings to live. A theory of human nature, drawing on the resources of the social as well as the physical sciences, could guide our reflections about what conduces to human flourishing. Psychoanalytic theory, for instance, could guide our views about whether a social world structured by certain values was mentally healthy or not. Williams proposes that if we did find that a social world promoted the best life or at least a flourishing life for human beings, this would justify the values embodied in that social world.

I have two reasons for placing this argument before you. First, I want you to consider Williams's very clear articulation of the idea of realism, and of what moral realism would be if we had it. The values of different cultures would represent their attempts to discern the moral part of reality, perhaps through the lens of some more perspectival concept. Second, I am interested in what the kind of objectivity Williams does suppose ethics might have has in common with scientific objectivity as he conceives it. This is the one piece of realism that I mentioned earlier as still haunting Williams's account, and it comes out clearly when Williams draws his conclusion. He argues that only one ethical belief might be objectively true in the ordinary, scientific, sense, namely, the belief that a certain sort of social world was the best one for human beings to live in. And then he says:

> other ethical beliefs would be true only in the oblique sense that they were the beliefs that would help us to find our way around in a social world which . . . was shown to be the best social world for human beings.[33]

Williams thinks that scientific beliefs are objective in the sense that they approximate as closely as possible to a representation of "the way the world

really is." Such beliefs help us find our way around in the natural world. So when Williams goes looking for some remnant of objectivity in ethics, he looks for the world that ethical beliefs would help us to find our way around in. Theoretical beliefs constitute a kind of map of the natural world, and ethical beliefs constitute a kind of map of the social world.[34] To be knowledge, a belief must help you find your way around in some world, he seems to think, and to do that, it must represent that world. It must describe some kind of reality.

IV. AGAINST THE MODEL OF APPLICATION

I think that this view of the relations between science and ethics represents a rather deep confusion about the difference between knowledge and action—something almost amounting to a failure to tell them apart. One way to articulate this admittedly difficult thought is by the metaphor I have just used. If to have knowledge is to have a map of the world, then to be able to act well is to be able to decide where to go and to follow the map in going there. The ability to act is something like the ability to *use* the map, and that ability cannot be given by *another map*. (Nor can it be given by having little normative flags added to the map of nature which mark out certain spots or certain routes as good. You still have to know how to use the map before the little normative flags can be of any use to you.) To put the same point another way, goodness in action cannot just be a matter of applying our knowledge of the good—not even a matter of applying our knowledge of what makes action itself good. This is because the ability to apply knowledge *presupposes* the ability to act.

Let me try to make this last point in a less metaphorical way. Suppose we agree that ethics is about what makes action good. For now I intend that phrase—"what makes action good"—to be neutral between consequentialist, virtue-theoretic, and deontological accounts of what makes an action good. Whatever their differences, the proponents of all of these accounts have to agree that one thing that makes an action good is conformity to the principle of instrumental reason, or what Kant calls the hypothetical imperative. An action that does not succeed in achieving its end is a failure, after all, but a movement that does not even succeed in aiming at an end doesn't succeed in being an action at all.

Now a realist account of the normativity of the instrumental principle is incoherent. For think how that account would have to work. The agent would have to recognize it, as some sort of eternal normative verity, that it is good to take the means to his ends. How is this verity supposed to motivate him? The obvious way to understand how facts motivate us is by means of a kind of extension of the instrumental principle itself. Philosophers have long acknowledged that the instrumental principle naturally extends to what is sometimes called "constitutive reasoning"—to use Williams's example, your end is an entertaining evening, and you choose dinner and a movie as what will constitute that end. The same line of thought extends the instrumental principle