1

THE PROBLEMS OF METAETHICS

1.1 What is metaethics?

Each year, as many as 300,000 girls undergo a procedure known as 'infibulation'; the World Health Organization calls it 'type III' genital cutting.[1] Extensive tissue is removed from their genitalia and the labia are stitched together so that after the procedure, nothing remains except a small opening for urine and menstrual blood. In many cases, infibulation is performed by a village midwife, with no anesthesia. There are no known health benefits of this procedure, but it can lead to bacterial infection, open sores, recurrent bladder and urinary tract infections, and increased risk of transmission of HIV. It also substantially raises the risks involved in childbirth: a 2006 study by the WHO found that infibulation raises the risk of the death of a child in pregnancy by 55 percent, raises the risk of cesarean section by 31 percent, and is associated with a 69 percent increase in the risk of postpartum hemorrhage, compared to women who have not undergone any genital cutting procedure.[2] Infants also need to be resuscitated in childbirth 66 percent more often, and their birth weight is on average 9 percent lower.[3]

Many people – you may be one of them – believe that it would be wrong to take one's pre-adolescent daughter to have this procedure performed, for no medical reason. Of those who think so, many believe that it is wrong to

allow others to have this procedure performed on their children, and as a result of views like these, the WHO has long sought to discourage procedures like this one. But some people disagree. They believe not only that it is wrong to prevent people from having procedures like this one performed on their daughters, but that it is permissible, or even a duty, to have this procedure performed on one's own daughter. Hundreds of thousands of parents choose to have the procedure performed on their own daughters every year, and millions more choose to have less drastic genital cutting procedures performed on their daughters.

What people disagree about, in this case, are *moral* questions: whether it is wrong or not to take one's daughter for infibulation, and whether it is wrong or not to *allow* people to take their daughters for infibulation. We are all familiar with moral questions, and we all have views about at least some of them. Sometimes the answers to moral questions seem easy or obvious – for example, most people find it obvious that killing an innocent person in cold blood in order to steal their DVD collection is wrong. You are not likely to find someone to disagree with you about this moral question, unless she is simply being disingenuous. But other times, even when the answer to a moral question seems obvious, you discover that other people disagree with you. For example, you discover that up to 300,000 girls each year are willingly subjected by their well-meaning parents to infibulation – parents who do not believe that what they are doing is wrong. Though you may find it obvious, they apparently do not. (Perhaps you are among those who do not find it obvious that infibulation is wrong – in that case, you disagree with the many people who do!)

Moments like this one tend to provoke a kind of existential paralysis known in philosophical circles as 'interest in metaethics', for reasons I'll explain in a moment. A first thought that you might have is: maybe taking one's daughter for infibulation is wrong *for us*, but all right *for them*. If you haven't had this thought, you are bound to have encountered someone who has. This is the idea that wrongness is *relative* – that actions are not simply right or wrong *simpliciter* – that is, all by themselves – but relative to a person, or relative to a cultural group, or relative to a time and place. Maybe, this idea claims, infibulation is wrong relative to our time and place, but all right, relative to the time and place of northern Sudan, where it is estimated that over 90 percent of women have undergone some variety of genital cutting procedure or other, and infibulation is very common.[4] The idea that wrongness is relative is a thesis about the *metaphysics* of morality. It is a thesis about what we are talking, thinking, or disagreeing *about*, when

we talk, think, or disagree about a moral question – that it is something relative, rather than absolute.

If you do not conclude that wrongness is relative, then next you are likely to wonder why you are so certain that infibulation is wrong, when other people are so equally certain that it is not wrong, and even that it is a duty. What makes you so sure, after all, other than that it seems obvious to you? But the opposite answer seems obvious to people with the opposing view. So what makes you think that what is obvious to you is a better guide to the truth than what seems obvious to them? If you are particularly reflective, you may add to these considerations the observations that given your upbringing and social circumstances, you were practically determined to find it obvious that infibulation is wrong, and that you would likely have found the reverse obvious, had you grown up in northern Sudan. These observations are bound to increase your puzzlement about how you really *know* that it is wrong, if your thinking that it is wrong is just a product of your upbringing and social circumstances. If you have ever wondered about any of these things, then you have been thinking about the *epistemology* of morality – the question of whether and how we do or can know the answers to moral questions.

Philosophers classify these questions about the metaphysics and epistemology of morality as belonging to the area of 'metaethics', so called because many people believe that the questions of metaethics are not questions *within* ethics – that is, they are not themselves moral questions – but are rather questions *about* moral questions – so they are 'meta' questions. Not everyone agrees with this characterization of what metaethics is about, but we have the name, nevertheless. Along with questions about the metaphysics and epistemology of morality, metaethics is concerned with questions about moral *thought* and questions about moral *language*. Since metaphysics, epistemology, the philosophy of mind (that is, of thought), and the philosophy of language are sometimes called the 'core areas' of philosophy, metaethics can therefore be characterized as what happens when we ask questions from the 'core areas' of philosophy about the subject matter of morality.

This book falls under the heading of 'metaethics' because it is primarily a book about moral *language*. But to appreciate the reasons why philosophers have been attracted to some of the theories about moral language that we will encounter in later chapters, it is important to see these questions about moral language as situated among related questions about moral reality (moral metaphysics), moral knowledge (moral epistemology), and moral thought (the philosophy of mind). And in fact, the questions about moral reality and moral knowledge are easier to understand.

1.2 The core questions (i): metaphysics and epistemology

We have already encountered one question about moral metaphysics, when you first entered your state of existential paralysis: it was our question as to whether wrongness is relative or absolute. But there are other, and more central, questions about moral metaphysics. The biggest is: what are moral questions *about*? The answer, of course, is that they are about morality – about what is right or wrong. But what kind of thing *is* that? Compare: if the question that we are interested in is whether sugar is soluble in water, there is more that we can say about what this is. For sugar to be soluble in water, after all, is for it to have the property that, very roughly, if you put some of it in water, then, other things being equal, it will dissolve. Now, that is not a very exciting answer to what questions about solubility are *about*, but it is an answer nonetheless. Metaethicists – people who spend their time thinking about metaethics – disagree a great deal over whether any answer at all can be given to the question, 'what are moral questions about?' other than the trivial one, 'they are about what is right or wrong', and if so, what kind of answer can be given.

One major concern that many people have had about the answer to this question is that they have noticed that morality is not a very scientific topic. Physicists and psychologists and biologists and chemists do not have very much to tell us about what is right and wrong – at least, not in virtue of the experiments or theories specific to their disciplines. Moreover, it is hard to see what sorts of experiments *would* be able to tell us the answers to our most important moral questions. People who believe that science is our best guide to what is true about the world therefore worry that we have a choice: either moral questions are really, ultimately about something that science can help to shed light on, or they are at best only about something spooky or unscientific.

The problem, however, is that it is hard to see how moral questions could possibly be about something that science can help to shed light on. Scientific investigation might reveal that babies do not, after all, die if they touch the clitoris during childbirth, as is reportedly believed by some in Burkina Faso.[5] Consequently, scientific investigation might *help us to settle* the question of whether it is wrong to not surgically remove a woman's clitoris before she is ready to give birth. But this kind of scientific investigation helps only if we know in advance that if a child will die if it touches the clitoris and if there is a high risk, if a woman with an intact clitoris gives birth, that the

infant touches it, then it must be wrong to not surgically remove a woman's clitoris before she gives birth. Since that is moral knowledge that we have to have *before* we do the scientific investigation, similar reasoning shows that not all moral knowledge can come from science – some of it has to come from somewhere else. And that, in turn, makes it seem to many people that moral questions can't really be about something that science can ultimately shed light on, and hence that if they are about anything at all, it is something spooky. Many metaethicists disagree with this conclusion, of course, but many also worry about it, to at least some degree or another.

We have also already encountered one of the main questions from moral epistemology: how do we *know* what we know about morality? It is important not to confuse this question, 'how do I know?' with the question, '*do* I know?' Just as you can see that when you press the button on a drinking fountain water comes out the spout, and be fully confident that it works while still wondering *how* it works, you can see that you know some things about what is right and wrong while still wondering *how* you know those things. The question of *how* we know what is right and wrong is the central question in moral epistemology.

Still, even though the questions, 'how do I know?' and '*do* I know?' are different questions, and we can ask the former while being confident that the answer to the latter is a robust 'yes', many philosophers have believed that the question of how we know what is right and wrong is particularly *hard*. And some people who are particularly impressed with how difficult this question is to answer have found themselves wondering '*do* I know?' after all.

To see why this question can seem to be particularly difficult, compare it to the question of how you know that there is a book in front of you. The sciences of optics, anatomy, neuroscience, and cognitive psychology have all contributed to helping us to understand how you could know such a thing. First, some light bounces off of the book and into your eyes. Then, it bounces off of your retina, where it encounters the photoreceptors that we know as rods and cones. From there, a signal passes through your optic nerve which corresponds to the pattern of stimulation of your rods and cones by different wavelengths and intensities of light, which is resolved by your visual cortex into a book-like shape. This is just part of the story, and a very sketchy part at that, but these and related sciences help to fill it in, so it turns out that scientists in fact know quite a lot about how you know that there is a book in front of you.

The problem is that there does not seem to be any similar such story that we can tell about how you know that it is wrong to take one's eight-year-old

daughter for infibulation. Unlike the book-like shape of the book in front of you, and the book-like pattern of inked text on its pages, the wrongness of taking one's eight-year-old daughter for infibulation is not, at least on the face of it, something that we can see. For that matter, it does not seem to be something that we can hear, or something that we can taste, touch, or smell, either. So how, exactly, do we find out about it? Recent research in anthropology, evolutionary biology, primatology, social psychology, brain imaging studies, and other fields adds to this worry, by telling us what *does* lead us, in general, to have the moral views that we do. Much of this research lends itself to the conclusion that we are evolutionarily destined to have certain views, that the reasons for which we think we hold them are really simply post hoc rationalizations, and that our moral thoughts are driven by our emotions, much more than by any kind of reasoning or reflective thought.[6]

When science gives us this kind of picture of where our moral views come from, it appears, at least initially, to contrast sharply with what science tells us about where your belief that there is a book in front of you comes from. And this is what makes it look hard to understand how, by processes like these, we ever manage to find out about anything – let alone about what is right or wrong. And that, in turn, leads some people who start with the question of *how* we know, to end up with the question of whether we know, after all.

1.3 The core questions (ii): mind and language

So far we have encountered one major question from moral metaphysics, and one major question from moral epistemology. We also need to introduce an important question from each of the philosophy of mind and the philosophy of language. These problems are somewhat more theory-laden than the problems we've just discussed from moral metaphysics and moral epistemology, so you may not have encountered them before. But they are questions that philosophers have become deeply puzzled about, and so it is important to at least see what they are. Fortunately, they are closely related to one another.

To see what these questions are, first notice that we have been talking about moral reality as something that moral questions are *about*. When you disagree with a man from northern Sudan about whether it is wrong to take one's daughter for infibulation, you are disagreeing *about* something. You are *talking* about this thing if you decide to go to Sudan to discuss it with him,

and you are *thinking* about this thing, when you ponder it at night. But just as we can wonder how you manage to *find out* about this thing – that is, how you *know* what is right or wrong – we can also wonder how you manage to *talk* and *think* about it. The question of how you manage to talk about it is a question from the philosophy of language, and the question of how you manage to think about it is a question in the philosophy of mind. These questions sound very similar, and in fact they are closely related.

To see how to become puzzled about these two questions, notice that many, many of the things that we manage to talk and think about are things with which we have become *acquainted* in some way – either directly or indirectly. For example, you have thoughts about the book in front of you (for example, you either think that it is pedantic, or that it is arcane, depending on your relative background in philosophy), and you are acquainted with it directly, by seeing it in front of you. You also have the ability to talk and have thoughts about Julius Caesar, one of which I have just gotten you to exercise. This is not because you are acquainted with him directly, but rather because you have heard about him from people who have heard about him from people … and so on, down to people who were directly acquainted with him. Hence, you are *indirectly* acquainted with Julius Caesar.

Acquaintance, either direct or indirect, or at least something like it, seems to play an important role in your ability to have thoughts about something. I can show this by giving you the ability to think about something that you have never before been able to think about. Unless you fall in a very small class of readers, you have never before had the ability to think or talk about my childhood dog, Chocolate Chip. But now you have the ability to both think and talk about Chocolate Chip; I have just given it to you, by giving you a name that I and my family used to refer to Chocolate Chip, while we were acquainted with him and his antics. So the ability to think or talk about someone or something doesn't come for free; it comes only under certain conditions, and the idea of direct and indirect acquaintance is one way of trying to spell out what those conditions involve, which gets the right results in the cases of this book, of Julius Caesar, and of Chocolate Chip.

But it is hard to see how rightness and wrongness are things that we can be acquainted with, either directly or indirectly in ways like these. If you can't see them, hear them, taste them, touch them, or smell them, then how does anyone become acquainted with them in the right sort of way? Thinking along these lines has made some philosophers who started with the question, '*how* does it happen that we manage to think and talk about what is right or wrong?' to become worried about the further question,

'*do* we manage to think or talk about what is right or wrong, *after all*?'. Certainly not all philosophers are worried about this question, but reflection on the difficulty of answering the 'how?' question has definitely made some people worry.

We've now encountered one major question from each of moral metaphysics, moral epistemology, the philosophy of moral language, and the philosophy of moral thought – each of the 'core areas' of philosophy as they apply to the subject matter of ethics. These questions were: what are moral questions *about*?; how do we *find out* about that?; how do we manage to *talk* about it?; and how do we manage to *think* about it? Similar questions from the 'core areas' of philosophy come up when philosophers think about many other topics, outside of ethics – for example, when philosophers think about mathematics, or causation, or material objects. So I will call them the *core questions*.

To understand one of the important reasons why at least some philosophers have been attracted to the kinds of theories about moral language that we will encounter in this book, it suffices to notice that each of the core questions contains a *presupposition* – an idea that it assumes to be true, such that it wouldn't make sense if that idea were not true. In all four cases, each question is based on the presupposition that moral questions are *about* something. If moral questions are not, really, *about* anything, then there is no real question to be answered as to *what* they are about. Likewise, if moral questions are not really about anything, then there is no real question to be answered as to how we find out about it. And similarly, there would be no real questions as to how we manage to talk about it or think about it. Hence, some philosophers who have thought that the core questions are difficult to answer in a satisfactory way have become attracted to the idea that the presupposition of these questions is in fact false – because moral questions aren't really *about* anything at all. If moral questions aren't really about anything, then we don't have to answer the difficult core questions – we can simply avoid them. And that is one of the main ideas whose consequences we are going to investigate in this book.

Because similar questions to our core questions come up when philosophers think about many different topics, there is an analogous move that is sometimes made in those other areas of philosophy. Just as some metaethicists conclude that moral questions are not really *about* anything, in order to avoid the presupposition of the core questions of metaethics, some philosophers of mathematics conclude that mathematical questions are not really *about* anything – or at least, not really about numbers – in order to avoid the

presupposition of the core questions of mathematics. And similarly, for other topics. Since this pattern happens in so many different places, in different areas of philosophy, it is useful to give it a name, and it is commonly called *nondescriptivism*, a usage I'll follow here. Hence, a *nondescriptivist* view about some subject matter holds that questions in that subject matter are not really *about* anything – at least in whatever sense is presupposed by the core questions. That is, they don't 'describe' anything – hence the name.

1.4 The motivation problem

There are many kinds of nondescriptivist theories in philosophy, about many different subject matters. And though the core questions constitute a *general* or *domain-neutral* motivation for nondescriptivism – which applies in many different areas of philosophy – some topics lend themselves to *special* or *domain-specific* motivations for nondescriptivism. For example, one of the most important problems in metaethics is called the *motivation problem*, and I'll explain what it is in just a moment. The motivation problem has led many philosophers to think that nondescriptivism about moral questions must be true. But the motivation problem doesn't arise in the philosophy of mathematics or in the philosophy of causation. So the reason it provides to believe in nondescriptivism is *special* or *topical*, in the sense that it trades on specific features of moral questions – things that seem to be true about moral questions but not about mathematical questions. That is what I will mean when I say that it is *domain-specific*, as opposed to the *domain-neutral* reasons for accepting nondescriptivism provided by the idea of avoiding the core questions by denying their presupposition.

The motivation problem starts with the following observation. Suppose that you and your friend have been discussing whether she ought to donate money to CARE, a highly rated international poverty-fighting organization.[7] She thinks not. Maybe she thinks that her money works more effectively to fight poverty if given to Oxfam,[8] or maybe she thinks that it is more important to donate to the political campaigns of the party she believes will make a larger difference than she can with her donation. Or maybe she simply thinks that it is her right to spend her money as she pleases, and prefers to spend it on soy lattes and sugar-free biscotti.

Whichever of these is the case, part of the point of engaging in this discussion with her is probably that you expect it to make a difference, if you convince her. Suppose, for example, you really do convince her that you are right, and that she ought to donate money to CARE. If the next thing

that happens is that a representative of CARE comes knocking on the door soliciting donations, you will expect that she will not be indifferent. Having decided that donating is what she ought to do, you will expect her to at least feel some motivation to donate. Before you convinced her, maybe she felt indifferent, but after you convince her, you expect her to feel indifferent no longer. If your friend really feels no such motivation, you are likely to wonder whether she was really just being insincere in agreeing with you, perhaps just hoping that you would get off her back.

So far, so obvious; all we've noted so far is that we typically expect people to feel at least some motivation to do what they think they ought to do. But I haven't gotten to the problem, yet. The problem is that this appears to make moral beliefs very different from other sorts of beliefs. If you are discussing with your friend whether the Alcove serves onion rings the size of doughnuts, for example, you *might* expect your friend to be motivated to eat lunch at the Alcove, but that depends on whether she has a hankering for outsized onion rings. If she goes in for such things, then she will be motivated to have lunch at the Alcove, if you suggest it, but if the thought of such onion rings disgusts her, then she will be motivated to go somewhere else. So in general, when you convince your friend of some ordinary, non-moral matter, there is nothing in particular that you expect that to motivate her to do (besides, of course, to admit that you are right about it) – what she will be motivated to do depends on what she wants or likes, what disgusts her or for what she hankers. But moral questions are different. If you convince your friend that she ought to donate to CARE, then there *is* something in particular that you expect this to motivate her to do, even if she has no special desire to do what she ought. You expect her to have some motivation to donate money to CARE.

The fact that moral beliefs seem, at least, to be different from non-moral beliefs in this way – that is, in terms of their motivational properties – is one that has struck philosophers as important for approximately as long as we have written records of anyone thinking about philosophy at all. The idea that moral beliefs *do* have a special connection to motivation that non-moral beliefs do not is called *motivational internalism*. Not everyone thinks that motivational internalism is true; a number of philosophers believe that there is something hasty and uncareful about the reasoning that I went through in the last two paragraphs in order to try to give you the idea of motivational internalism. They believe that if you think more carefully about how moral beliefs motivate, you will see that they are not, in fact, all that different from non-moral beliefs. These philosophers accept what is known as *motivational externalism*.

Still, even though there is much disagreement about the truth of motivational internalism – that is, about whether moral beliefs have a special connection to motivation that non-moral beliefs do not – people who *do* think that it is true are led next to wonder *why* it is true. So even if you, like the motivational externalists, doubt that there is any very deep difference underlying the two examples that I just gave, suppose for just a moment that you really were convinced that there is such a deep difference between moral beliefs and non-moral beliefs. So what we want to know is: *why*?

The most influential answer to this question has been that moral beliefs have a special connection to motivation that non-moral beliefs do not because moral beliefs are a *different kind of mental state* from non-moral beliefs. Proponents of this answer often pair it with a picture of *how* these two kinds of mental states differ. According to the picture, non-moral beliefs are *about* something. They are like maps, which tell us the lay of the land, where 'the land' is a metaphor for what they are about. Whereas moral beliefs, according to this picture, are not maps to anything. They are more like *goals* we have, about which destination on the map we want to reach. According to this picture, ordinary, non-moral beliefs do not correspond to any particular thing that they motivate someone to do, because even once we know where we are on the map, what we do next depends on what destination we are trying to reach. Whereas according to the picture, moral beliefs do correspond to a particular thing that they motivate someone to do, because having a moral belief *is* a matter of having some particular destination as your goal.

The picture that I have just sketched, according to which some mental states are like maps of the world, and other mental states are like goals about which destination we are headed for, is sometimes called the *Humean Theory of Motivation*. It is called that because, in the eighteenth century, David Hume drew this picture in a very compelling way.[9] A different metaphor sometimes used to describe what is essentially the same picture, is that of *direction of fit*.[10] According to this metaphor, some mental states 'try' to match to the world, and are unsuccessful if they do not change in order to match the world. Whereas other mental states 'try' to get the world to match themselves, and are unsuccessful if they do not get the world to change in order to match them. The difference between these two sorts of mental states is therefore like the difference between the notepad of a detective who follows a grocery shopper, trying to keep track of what goes into her cart, and the shopping list of that shopper. Both lists try to match what is in the cart, but the detective tries to do it by changing his list, and the shopper tries to do it by changing what is in the cart.

So put in terms of the Humean Theory of Motivation, the most influential answer to the question of *why* moral beliefs have a special tie to motivation that non-moral beliefs do not, is that non-moral beliefs are like the detective's list, and have what is called 'mind-to-world' direction of fit, because they try to match – to *map out* – how the world is. Whereas moral beliefs have a special tie to motivation, because they are like the shopper's list, and have what is called 'world-to-mind' direction of fit, because they try to get the world to match themselves. These are just metaphors, but they are helpful metaphors, because they give us a picture that different theories fill out in different, more concrete ways.

1.5 Noncognitivism in ethics

We now know enough to be able to say what the motivation problem has to do with nondescriptivism, the view that we encountered in the last section. If moral beliefs contrast with non-moral beliefs in not being maps to anything – in there being nothing about the world that they try to match – then in a natural sense they are not *about* anything. After all, intuitively the thing that a belief is about, is the very thing that it tries to match, and is unsuccessful if it does not match. So the influential answer to the question of why motivational internalism is true turns out to be the same as the strategy for avoiding the core questions, by denying their presupposition. Both involve the idea that, in at least some sense, moral beliefs are not really *about* anything.

This means that the strategy for avoiding the core questions and the influential answer to why motivational internalism is true lead to the same place. They both lead to *nondescriptivism*. Because of its connection to the motivational problem, nondescriptivism about morality often goes by another name. It is called *noncognitivism*. If 'cognitive' means 'having to do with belief', the idea is that moral thoughts are not of the same kind as ordinary beliefs – at least, ordinary *non*-moral beliefs. So they are in that sense 'noncognitive', and hence the name for the view.

Philosophers have offered many different definitions for the term 'noncognitivism' over the last seven decades or more. Some have defined it as the view that moral sentences cannot be true or false. Some have defined it as the view that there is no such thing as a moral belief. Some have defined it as a special version of *expressivism*, a particular theory that we will encounter in Chapter 4. So when you read other books and articles about this topic, you will see that the word is used differently by different people. But one thing

that people generally agree on, is *who counts* as a noncognitivist. It is agreed by most philosophers that the theories of A.J. Ayer, Charles Stevenson, R.M. Hare, Simon Blackburn, and Allan Gibbard all count as part of the 'noncognitivist tradition'. Since we don't have a better name for this tradition, that is how I will use the word 'noncognitivism' in this book.

We will begin to encounter the views of these theorists in Chapter 2, where we will see that as a matter of fact, there is not very much that they can all agree on. But they do agree on this much: that we don't need to worry about the core questions, because they have a false presupposition, and that as a result, we can explain why moral thought and language are more intimately connected to motivation than non-moral thought and language. So I will understand noncognitivism in ethics to be a family of nondescriptivist theories about morality which try to explain a special connection between moral thought and motivation.

And so I can now (at last!) say what this book is about: it is about the problems and prospects facing noncognitivist theories. So far, we've seen in this chapter that noncognitivist theories can be perceived to have certain advantages in the core areas of philosophy, and that they can apparently provide a neat explanation of motivational internalism, which is at least very plausible as an observation about moral thought. These are the main *attractions* of noncognitivism – the reasons why a variety of philosophers are tempted to think that it is true. Now, you may not find these particular features attractive. For example, you may think that there *are* satisfactory answers to what moral questions are about, and as to how we find out about and manage to talk and think about what they are about.. And you may think that there was something misleading about my examples in which you convince your friend that she ought to donate, as compared to the case in which you convince her that there are doughnut-sized onion rings at the Alcove, and hence you might think that motivational internalism is a big exaggeration. If you think those things, then noncognitivism will not be attractive *to you*. But many people have become convinced that the core questions are particularly difficult in the case of morality, and that motivational internalism *is* true. So noncognitivism is attractive to *them*. (And you may, in fact, be one of them.)

So how, then, do we figure out who is right? If we disagree about how difficult the core questions are to answer satisfactorily, and about whether motivational internalism is true, then how do we decide whether noncognitivism is true? Well, there are essentially two ways that we can go. On the first strategy, we spend our efforts trying to figure out what

the right answers are to the core questions, assuming that noncognitivism is false. If we *can* answer these questions in a satisfactory way (even if they are quite difficult), then we learn something: we learn that at the very least, these questions are not *too* hard to answer. And that removes one kind of reason – the reason based on the strategy for avoiding the core questions – to think that noncognitivism is true. Similarly, as part of this strategy we can spend our efforts trying to determine whether motivational internalism is true – that is, just what the nature of the connection is between moral belief and motivation, and whether we can explain it by any less drastic measures than those proposed by the noncognitivist. If we can, then again, that removes one of the reasons to think that noncognitivism is true.

So the first strategy is *indirect*. It says: find out the answers to everything else in metaethics, and if we can reach satisfactory answers, then there may be no reason to think that noncognitivism is true. But the second strategy for trying to figure out whether noncognitivism is true is more direct. Setting aside whether it is well motivated, on the second strategy we spend our efforts trying to figure out what things would be like if noncognitivism were true, and checking to see whether things could really be that way. In order to explore this strategy, we need to know more of what noncognitivist theories are like. We need to begin to develop the details of these theories, and to see whether they have testable predictions.

Nearly any other book about metaethics will tell you some of the things that you need to know in order to work on the first strategy. In this book, we will take the second strategy. Having explained some of the reasons why people are tempted to think that noncognitivism is true, I am now going to set those reasons aside for the remainder of the book, except when we are trying to figure out whether a given noncognitivist theory really achieves the main purported advantages of noncognitivism, after all. Instead of trying to figure out whether noncognitivism is a *tempting* view, we will spend our time trying to figure out whether it *could* be true.

In particular, as we'll see in Chapter 2 when we learn more about what individual noncognitivist theories look like, it will turn out that just as noncognitivism is thought to *solve* various problems from the core areas of philosophy, noncognitivist theories also *face* important problems in each of the core areas of philosophy. The general reason why noncognitivism faces these problems is simple: if moral questions are not *about* anything, and we are not talking about anything when we discuss them – any special subject matter of morality – then what *are* we doing?

We do, after all, have moral thoughts. You and I both think, for example, that it is wrong to kill an innocent person in order to steal their DVD collection. So if this thought is not *about* anything in particular – the wrongness of this action – then what exactly *is* involved in having this thought? Similarly, we do use moral language – we say to each other sentences like, 'Infibulation is wrong' or 'It's not wrong to take one's daughter for infibulation'. So if these sentences are not *about* anything in particular – the wrongness or lack of it attaching to a certain surgical procedure – then what do they *mean*, and what use are they? Noncognitivist theories in ethics face a special burden to answer these questions – to tell us what kind of thoughts moral thoughts are, and what kind of *meaning* moral sentences have, if they aren't about anything in particular. It is in answering these questions that noncognitivist theories run into their problems.

So in later chapters, we will look at problems faced by noncognitivist theories in the philosophy of mind (Chapter 5), in the philosophy of language (Chapters 3, 6, and 7), in metaphysics (Chapter 8), and in epistemology (Chapter 9). Along the way, we'll gradually learn about what different kinds of noncognitivist theory are like (Chapters 2, 4, and 10). Unless noncognitivists can solve these problems that face their theories, noncognitivism *can't* be true, even if it could help us to solve many philosophical problems if it *were* true.

Chapter summary

In this chapter we introduced the *domain-neutral* core questions of metaphysics, epistemology, the philosophy of language, and the philosophy of mind, as well as the *domain-specific* motivation problem from metaethics. We introduced noncognitivism as a kind of *nondescriptivist* theory about what kind of meaning moral words have, if they are not about anything. And we previewed what to expect in the remainder of the book.

Further reading

An alternative take on the central problems of metaethics can be found in Darwall, Gibbard, and Railton (1997); Miller (2003) provides a fairly typical introduction to the field and its various topics. Both are written at a somewhat more advanced level than this chapter. Urmson (1968, chapter 2) is also recommended. I follow Urmson in distinguishing between domain-neutral motivations from the core areas, and the domain-specific problem about motivation.

Exercises

1 E *Comprehension*: In section 2 we distinguished between the question of *whether* we have moral knowledge and the question of *how* we get that knowledge, and in section 3 we distinguished between the question of whether we manage to talk and think about moral questions and how we manage to do so. Provide examples of two other topics about which you might wonder both whether something is true, and how it is true. Choose one example about which you wonder how something is true, without that making you doubt that it is true. Choose another example, for which wondering how something could be true makes you wonder whether it really is true.

2 E *Qualifications*: In the main text, I characterized noncognitivist views as holding that moral sentences are not really about *anything*. This characterization was somewhat sloppy. Consider the following four sentences:

1 Infibulation is wrong.
2 Infibulation is common.
3 Infibulation is both common and wrong.
4 Aww, infibulation!?

For each sentence, list the things that sentence is intuitively about. Can you list the same number of things for each sentence, or do you get different numbers for different sentences? How do these numbers compare?

3 E *Branching out*: The Open Question argument. Rank the following four questions from most to least interesting or 'live' feeling (ties are allowed):

5 I know it's wrong, but is it wrong?
6 I know it's wrong, but is it harmful?
7 I know it's harmful, but is it wrong?
8 I know it's harmful, but is it harmful?

Do your rankings change if you substitute 'against God's will' for 'harmful'? If you substitute 'untraditional' for 'harmful'? Can you think of anything that you could substitute for 'harmful' that would make your rankings change?

4 E *Extensions*: Rank the following four questions, as you did in the previous exercise:

9 I know it's a triangle, but is it a triangle?
10 I know it's a triangle, but is it a shape with three straight sides?
11 I know it's a shape with three straight sides, but is it a triangle?
12 I know it's a shape with three straight sides, but is it a shape with three straight sides?

Do these questions seem the same or different? 'Shape with three straight sides' seems like a pretty good definition of 'triangle'. Could there be a definition for 'wrong'?

5 M *Extensions*: We noted in the main text that the motivation problem is *domain-specific*, meaning that it does not come up in every area of philosophy. In order to test whether this is true, try to construct an example like the one that we used to introduce motivational internalism for each of two other subject matters: mathematics and beauty. Try to construct a case of a mathematical question – for example, from arithmetic, algebra, or calculus – such that there is something that you expect your friend to be motivated to do, if you convince her of the answer to that question, and not just because of what she wants or likes. And try to construct a similar example for a question about what is beautiful. Is one of these harder to do than the other? How do they compare to the moral case?

6 M *Branching out*: It is natural to wonder whether there are domain-specific problems that motivate nondescriptivism about other topics. For a good example, see Gibbard (1981), who offers a simple argument for nondescriptivism about *conditionals* ('if ... then' sentences) that doesn't apply to other topics. Another good example is the case of truth. Why might the following sentence make someone puzzled about how there could be anything we are saying about the world when we say that a sentence is true?

Liar: The sentence named 'Liar' on page 17 of *Noncognitivism in Ethics* is not true.

Morals

2 In addition to being about wrongness, sentence 1 is about infibulation with excision. Noncognitivist views don't deny that sentence 1 is about infibulation with excision; just that it is about something else – wrongness. So

it is sloppy to say that according to noncognitivist views, moral sentences are not about anything. What is important about noncognitivist views, is that a sentence is not about *more* things, in virtue of having a moral word like 'wrong' in it, because 'wrong' does not *contribute* to what the sentence is about.

3 One reason some philosophers have been skeptical that there is any answer to the question of what it is to be wrong is because of an argument advanced by G.E. Moore known as the *Open Question argument*. Moore believed that no matter what is substituted for 'harmful', question 7 feels 'open' in a way that questions 5 and 8 do not. This led Moore to believe that they are different questions, and since they are different questions, to believe that whatever is substituted for 'wrong' can't be the answer to what it is to be wrong. Some philosophers believe that even if these sentences all ask the same question, the use of different words to ask it can make 6 and 7 feel 'open', even though 5 and 8 do not. This view predicts that 6 and 7 should feel *equally* 'open'. Did you say that they are equally live questions? Does it matter what we substitute for 'harmful'?

4 One good way to test Moore's Open Question argument to see whether it is a good argument is to check and see whether true definitions feel 'open'. Another good kind of definition to test is the fact that water is H_2O. Do the questions in exercise 3 feel 'open' in a way that the analogous questions with 'water' substituted for 'wrong' and 'H_2O' substituted for 'harmful' do not? Or do they feel the same? Some philosophers believe that it is just the same, and some believe that there is something special and different going on in the moral case, which makes it harder to say what 'wrong' is about than what 'triangle' or 'water' are about.

References

Anscombe, Elizabeth (1957). *Intention*. Oxford: Basil Blackwell.

Darwall, Stephen, Allan Gibbard, and Peter Railton, eds. (1997). *Moral Discourse and Practice: Some Philosophical Approaches*. Oxford: Oxford University Press.

de Waal, Frans (1996). *Good Natured: The Origins of Right and Wrong in Primates and Other Animals*. Cambridge, MA: Harvard University Press.

Greene, J.D., and J. Haidt (2002). 'How (and Where) Does Moral Judgment Work?' *Trends in Cognitive Sciences* 6: 517–23.

Greene, J.D., R.B. Sommerville, L.E. Nystrom, J.M. Darley, and J.D. Cohen (2001). 'An fMRI Investigation of Emotional Engagement in Moral Judgment.' *Science* 293: 2105–8.

Haidt, Jonathan (2001). 'The Emotional Dog and Its Rational Tail: A Social Intuitionist Approach to Moral Judgment.' *Psychological Review* 108: 814–34.

Joyce, Richard (2006). *The Evolution of Morality*. Cambridge, MA: MIT Press.

Miller, Alexander (2003). *An Introduction to Contemporary Metaethics*. Cambridge: Polity.

Nichols, Shaun (2004). *Sentimental Rules*. Oxford: Oxford University Press.

Smith, Michael (1994a). *The Moral Problem*. Oxford: Basil Blackwell.

Urmson, J.O. (1968). *The Emotive Theory of Ethics*. New York: Oxford University Press.